

Working With Data

Programming in R for Data Science

Anders Stockmarr, Kasper Kristensen, Anders Nielsen

Data frames; adding and removing columns

- ▶ Look at the data frame dat:

```
> dat <- data.frame(x=LETTERS[1:3], y=1:3)
```

```
> dat
```

```
  x y
```

```
1 A 1
```

```
2 B 2
```

```
3 C 3
```

```
> dat[,1]
```

```
[1] A B C
```

```
Levels: A B C
```

```
> dat$x
```

```
[1] A B C
```

```
Levels: A B C
```

- ▶ It is simple to add or remove a column:

```
> dat$z <- dat$y^2
```

```
> dat$name <- c("Cat", "Vic", "Osc")
```

```
> dat$y<-NULL
```

```
> dat
```

```
  x z name
```

```
1 A 1  Cat
```

```
2 B 4  Vic
```

```
3 C 9  Osc
```

Merging data frames

- If we have two data sets:

```
> dat1
```

	name	age
1	Cat	9
2	Vic	7
3	Osc	4

```
> dat2
```

	names	gender
1	Vic	Male
2	Cat	Female
3	Osc	Male

- Then we can merge that information into one data set by:

```
> dat <- merge(dat1, dat2, by.x="name", by.y="names")
```

```
> dat
```

	name	age	gender
1	Cat	9	Female
2	Osc	4	Male
3	Vic	7	Male

Getting dimension and column info

Given the dataset:

```
> df
```

	name	age	gender
1	Cat	9	Female
2	Osc	4	Male
3	Vic	7	Male

► names

```
> names(df)
```

```
[1] "name"    "age"     "gender"
```

► class

```
> class(df$name)
```

```
[1] "factor"
```

```
> class(df$age)
```

```
[1] "numeric"
```

► dim

```
> dim(df)
```

```
[1] 3 3
```

```
> nrow(df)
```

```
[1] 3
```

```
> ncol(df)
```

```
[1] 3
```

Object structure

- Get overview of the object structure:

```
> str(df)
```

```
'data.frame':      3 obs. of  3 variables:
 $ name  : Factor w/ 3 levels "Cat","Osc","Vic": 1 2 3
 $ age   : num  9 4 7
 $ gender: Factor w/ 2 levels "Female","Male": 1 2 2
```

First and last rows

- ▶ First rows of a data frame:

```
> head(airquality, 3)
```

	Ozone	Solar.R	Wind	Temp	Month	Day
1	41	190	7.4	67	5	1
2	36	118	8.0	72	5	2
3	12	149	12.6	74	5	3

- ▶ Last rows of a data frame:

```
> tail(airquality, 3)
```

	Ozone	Solar.R	Wind	Temp	Month	Day
151	14	191	14.3	75	9	28
152	18	131	8.0	76	9	29
153	20	223	11.5	68	9	30

The subset() function

Let's look at the `airquality` data again:

```
> head(airquality, 3)
```

	Ozone	Solar.R	Wind	Temp	Month	Day
1	41	190	7.4	67	5	1
2	36	118	8.0	72	5	2
3	12	149	12.6	74	5	3

- ▶ Logical indexing applies to data frames:

```
> dataA <- airquality[airquality$Temp>80, c("Ozone", "Temp")]
```

- ▶ But a neat function is built in for making subsets of data

```
> dataA <- subset(airquality, Temp > 80, select = c(Ozone, Temp))
```

```
> dataB <- subset(airquality, Day == 1, select = -Temp)
```

```
> dataC <- subset(airquality, select = Ozone:Wind)
```

The `summary()` function

- ▶ The **`summary()`** function gives you a range of statistics

```
> summary(airquality$Wind)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.700	7.400	9.700	9.958	11.500	20.700

that you could alternatively obtain using the **R** functions **`min()`**, **`max()`**, **`mean()`**, **`median()`**, **`quantile()`**.

- ▶ The summary of a data frame gives the summary of each column:

```
> summary(airquality)
```

Ozone	Solar.R	Wind	Temp	Month	Day
Min. : 1.00	Min. : 7.0	Min. : 1.700	Min. :56.00	Min. :5.000	Min. : 1.0
1st Qu.: 18.00	1st Qu.:115.8	1st Qu.: 7.400	1st Qu.:72.00	1st Qu.:6.000	1st Qu.: 8.0
Median : 31.50	Median :205.0	Median : 9.700	Median :79.00	Median :7.000	Median :16.0
Mean : 42.13	Mean :185.9	Mean : 9.958	Mean :77.88	Mean :6.993	Mean :15.8
3rd Qu.: 63.25	3rd Qu.:258.8	3rd Qu.:11.500	3rd Qu.:85.00	3rd Qu.:8.000	3rd Qu.:23.0
Max. :168.00	Max. :334.0	Max. :20.700	Max. :97.00	Max. :9.000	Max. :31.0
NA's :37	NA's :7				

Missing values

- ▶ **R** uses the special value **NA** to code missing values.
- ▶ The result of arithmetic involving NAs becomes NA as well:

```
> colMeans(airquality)
```

Ozone	Solar.R	Wind	Temp	Month	Day
NA	NA	9.957516	77.882353	6.993464	15.803922

- ▶ This also holds for the comparison operator:

```
> NA == NA
```

```
[1] NA
```

Missing values

- ▶ We need a special function **is.na** to filter out NAs:

```
> is.na(NA)
[1] TRUE
```

- ▶ To get rid of NAs in a column we can use

```
> s <- subset(airquality, !is.na(Ozone) )
> colMeans(s)
```

Ozone	Solar.R	Wind	Temp	Month	Day
42.129310	NA	9.862069	77.870690	7.198276	15.534483

- ▶ Note that the argument **na.rm=TRUE** can be passed to most summary functions e.g. **sum()**, **mean()**, **sd()**:

```
> mean(airquality$Ozone, na.rm=TRUE)
[1] 42.12931
```

Text manipulation

Text mining/text analytics is a relatively new and evolving science; online text mining has evolved in this century.

- ▶ Here are some examples:
 - ▶ Tweets.
 - ▶ Google Flu Trends (now terminated).
 - ▶ Questionnaires with freeform text.
 - ▶ Extraction of quantitative information from messy text sources.
- ▶ Text mining can be done in R.
- ▶ **R** has some simple tools that can be readily used for text searches.

Text search and replace

- Overview:

Function	Description
grep() grepl()	Text pattern search
sub()	Text pattern search and replace first occurrence
gsub()	Text pattern search and replace all occurrences

- Note common arguments:

- By default pattern is a *regular expression*: **fixed = FALSE**.
- By default matching is *case sensitive*: **ignore.case = FALSE**.

- Lets try some examples...

Text search example

- ▶ Some example text

```
> txt <- c("Hello, my",  
+         "name is",  
+         "anders."  
+         )
```

- ▶ Search for string 'name' returning indices:

```
> grep("name", txt)
```

```
[1] 2
```

or logical indices

```
> grepl("name", txt)
```

```
[1] FALSE TRUE FALSE
```

- ▶ Replace 'anders' with 'Anders':

```
> sub("anders", "Anders", txt)
```

```
[1] "Hello, my" "name is"   "Anders."
```

Text manipulation - Regular expressions

- ▶ Messy dataset: Number of fruits eaten by person.

```
> df
```

	person.ID	fruit
1	1	apple: 3 Orange : 9 banana:2
2	2	Orange:1 Apple: 3 banana: 10
3	3	banana: 3 Apple: 3 Orange : 04

- ▶ We want the number of oranges eaten for each person. How to get this information automatically?
- ▶ Answer: Construct a regular expression to extract it...

Regular expressions

- ▶ Regular expression to find the orange count:

`.*` A sequence of arbitrary characters (possibly empty) followed by
`orange` followed by
`[:]*` a sequence of whitespace OR colon followed by
`([0-9]*)` a sequence of digits followed by
`.*` a sequence of arbitrary characters (possibly empty).
▶ Parenthesis around the digit pattern stores this match. The match is accessed through a so-called *back reference* `"\1"`.

- ▶ Now try it:

```
> pattern <- ".*orange[ :]*([0-9]*).*"  
> sub(pattern, "\\1", df$fruit, ignore.case=TRUE)  
  
[1] "9"  "1"  "04"
```

- ▶ Use `as.numeric()` to convert result to a number.

Date and time objects

- ▶ **R** has classes to handle dates **?as.Date** and dates with time information **?as.POSIXct**.

- ▶ You can construct these manually from characters:

```
> as.Date("2016-03-10")
```

```
[1] "2016-03-10"
```

```
> as.POSIXct("2016-03-10 13:53:38 CET")
```

```
[1] "2016-03-10 13:53:38 CET"
```

- ▶ But if you are lucky your SQL server already uses a suitable datetime format that will be properly imported by **R** :

```
> conn <- odbcDriverConnect(connStr)
```

```
> df <- sqlQuery(conn, "SELECT TOP 2000 * FROM bi.sentiment")
```

```
> class(df$date)
```

```
[1] "POSIXct" "POSIXt"
```


Working with dates

- ▶ Many useful methods exist for dates (see `methods(class="Date")`):
- ▶ You can do math with dates / datetimes.

```
> mean(df$Date)
```

```
[1] "2014-01-13 16:40:04 CET"
```

```
> mean(df$Date+10)
```

```
[1] "2014-01-13 16:40:14 CET"
```

- ▶ However, note the unit difference:

```
> mean(as.Date(df$Date))
```

```
[1] "2014-01-12"
```

```
> mean(as.Date(df$Date)+10)
```

```
[1] "2014-01-22"
```

- ▶ Some useful conversion methods:

```
> table weekdays(df$Date)
```

```
l rdag  onsdag
```

```
819    1181
```

```
> table(months(df$Date))
```

```
februar  januar
```

```
819    1181
```